

Sheng Xu

In [1]:

```
from LRGradientDescent import LogisticRegressionGradientDescent as LRGD
import numpy as np
from scipy.special import logsumexp
from scipy.special import expit as sigm #sigmoid function

import os
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

In [2]:

```
## Toy problem
#
# Logistic regression should be able to perfectly predict all 10 examples
# five examples have x values within (-2, -1) and are labeled 0
# five examples have x values within (+1, +2) and are labeled 1
N = 10
x_NF = np.hstack([np.linspace(-2, -1, 5), np.linspace(1, 2, 5)][:, np.newaxis])
y_N = np.hstack([np.zeros(5), 1.0 * np.ones(5)])

lr = LRGD(
    alpha=0.1, step_size=0.1, init_w_recipe='zeros')

# Prepare features by inserting column of all 1
xbias_NG = lr.insert_final_col_of_all_ones(x_NF)
```

1a

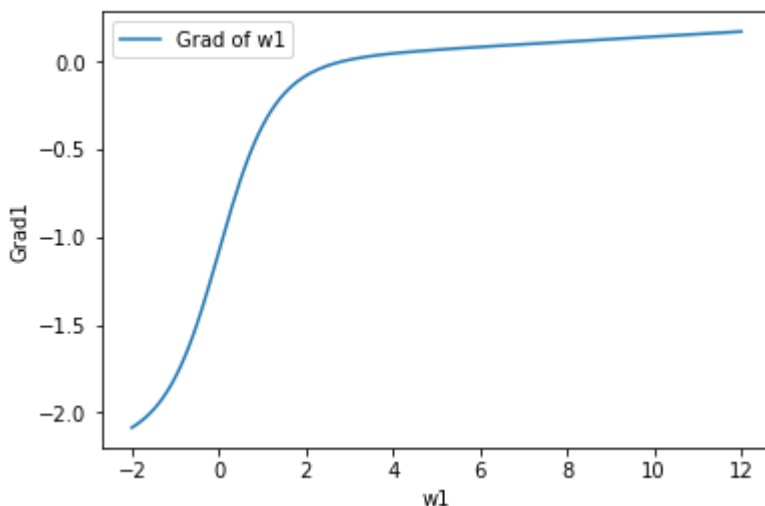
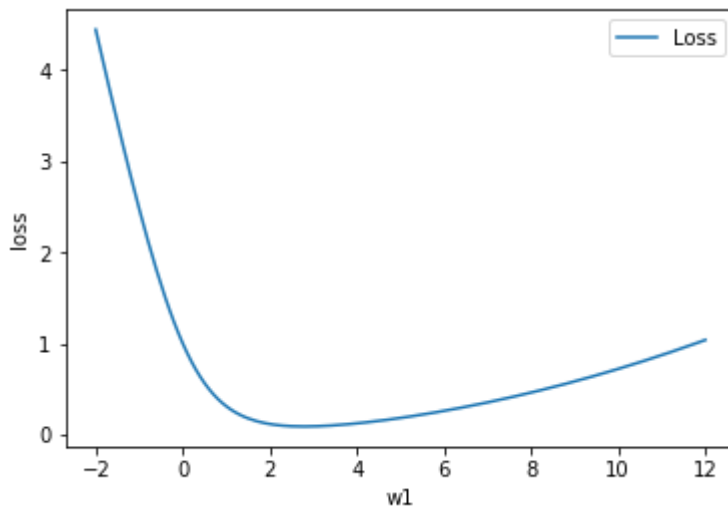
In [3]:

```
loss=[]; grad=[];
arr=np.linspace(-2, 12, 141)
for i in arr:
    w=np.array([i,0])
    loss.append(lr.calc_loss(w, xbias_NG, y_N))
    grad.append(lr.calc_grad(w, xbias_NG, y_N)[0])
#print (loss, grad)
id_min_cost=np.argmin(loss)
```

In [4]:

```
plt.plot(arr, loss, label='Loss')
plt.xlabel('w1');
plt.ylabel('loss');
plt.legend();
plt.show();

plt.plot(arr, grad, label='Grad of w1')
plt.xlabel('w1');
plt.ylabel('Grad1');
plt.legend();
plt.show();
```



Yes.

For loss:

We know that when $w1$ is less than 0, the loss is big because the value of \log_loss is big, because $w * x$ is negative, which leads to wrong classification.

On the other hand, when $w1$ is positive, the loss is bigger because the $l2_penalty$ goes bigger as $w1$ is bigger.

For Gradient:

From the information of loss, we know: the partial derivative of $w1$ is negative when $w1 < 0$, it will reach 0 somewhere positive. Then, it goes positive when $w1$ grows bigger.

The minimum is somewhere between 2 and 3, close to 3.

Here, my estimation is 2.8. (See below)

In [5]:

```
print("best w1 for LR with 1 feature and 0 bias: %.3f" % arr[id_min_cost])
```

best w1 for LR with 1 feature and 0 bias: 2.800

In [6]:

```
lr.fit(x_NF, y_N)
```

Initializing w_G with 2 features using recipe: zeros

Running up to 10000 iters of gradient descent with step_size 0.1

iter	0/10000	loss	1.000000	avg_L1_norm_grad	0.541011
w[0]	0.000	bias	0.000		
iter	1/10000	loss	0.888015	avg_L1_norm_grad	0.494016
w[0]	0.108	bias	0.000		
iter	2/10000	loss	0.794586	avg_L1_norm_grad	0.451748
w[0]	0.207	bias	0.000		
iter	3/10000	loss	0.716373	avg_L1_norm_grad	0.414112
w[0]	0.297	bias	0.000		
iter	4/10000	loss	0.650555	avg_L1_norm_grad	0.380787
w[0]	0.380	bias	0.000		
iter	5/10000	loss	0.594813	avg_L1_norm_grad	0.351344
w[0]	0.456	bias	0.000		
iter	6/10000	loss	0.547278	avg_L1_norm_grad	0.325330
w[0]	0.527	bias	0.000		
iter	7/10000	loss	0.506451	avg_L1_norm_grad	0.302308
w[0]	0.592	bias	0.000		
iter	8/10000	loss	0.471140	avg_L1_norm_grad	0.281878

1b

In [7]:

```
print(" Result for LR with 1 feature and 0 bias: ", lr.trace_w[-1])
```

Result for LR with 1 feature and 0 bias: [2.78265835e+00 1.02172570e-17]

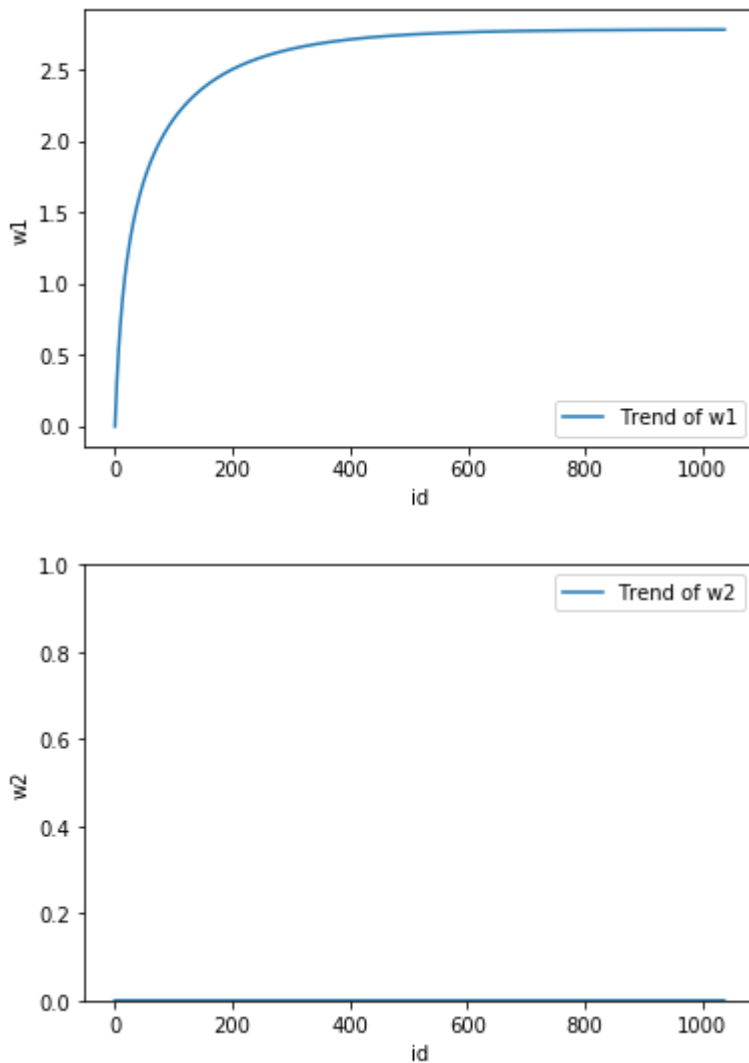
In [8]:

```
matrix=np.matrix(lr.trace_w).T
w1=np.asarray(matrix[0])[-1]
w2=np.asarray(matrix[1])[-1]
idx=np.linspace(0, w1.size-1,w1.size)
```

In [9]:

```
## Draw Picture
plt.plot(id, w1, label='Trend of w1')
plt.xlabel('id');
plt.ylabel('w1');
plt.legend();
plt.show();

plt.plot(id, w2, label='Trend of w2')
plt.xlabel('id');
plt.ylabel('w2');
plt.ylim([0.0, 1.0]);
plt.legend();
plt.show();
```



Yes.

w_1 approaches 2.78. The converging speed is fast at first, then becomes really slow before w_1 reaches the final result.

w_2 stays at nearly 0 because it should be 0 from the symmetry of the question.

1c

In [10]:

```
lr2 = LRGD(
    alpha=0.1, step_size=0.1, init_w_recipe='uniform_-1_to_1')
lr2.fit(x_NF, y_N)
```

Initializing w_G with 2 features using recipe: uniform_-1_to_1

Running up to 10000 iters of gradient descent with step_size 0.1

iter	0/10000	loss	0.932814	avg_L1_norm_grad	0.579579
w[0]	0.098	bias	0.430		
iter	1/10000	loss	0.834339	avg_L1_norm_grad	0.534178
w[0]	0.198	bias	0.415		
iter	2/10000	loss	0.751280	avg_L1_norm_grad	0.492499
w[0]	0.290	bias	0.400		
iter	3/10000	loss	0.681012	avg_L1_norm_grad	0.454653
w[0]	0.374	bias	0.385		
iter	4/10000	loss	0.621302	avg_L1_norm_grad	0.420533
w[0]	0.452	bias	0.372		
iter	5/10000	loss	0.570293	avg_L1_norm_grad	0.389904
w[0]	0.523	bias	0.360		
iter	6/10000	loss	0.526460	avg_L1_norm_grad	0.362461
w[0]	0.590	bias	0.348		
iter	7/10000	loss	0.488565	avg_L1_norm_grad	0.337881
w[0]	0.651	bias	0.337		
iter	8/10000	loss	0.455603	avg_L1_norm_grad	0.315844

In [11]:

```
print(" Result for LR with 1 feature and 0 bias: ", lr2.trace_w[-1])
```

Result for LR with 1 feature and 0 bias: [2.78266077e+00 9.11289218e-04]

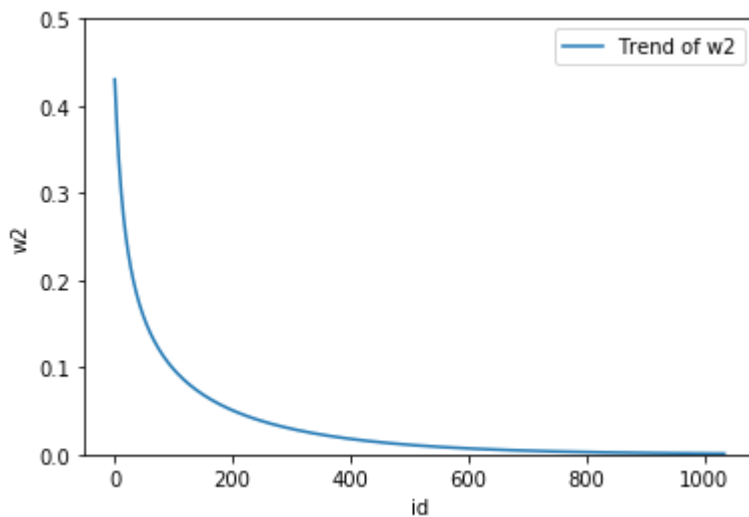
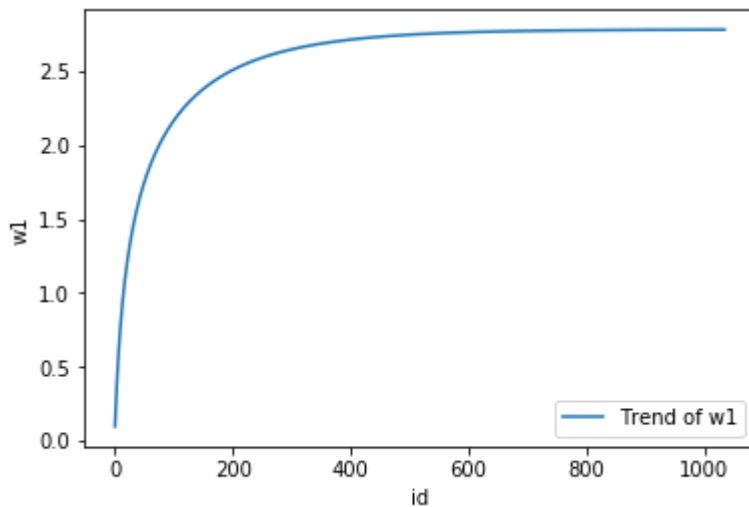
In [12]:

```
matrix=np.matrix(lr2.trace_w).T
w1=np.asarray(matrix[0])[-1]
w2=np.asarray(matrix[1])[-1]
idx=np.linspace(0, w1.size-1,w1.size)
```

In [14]:

```
## Draw Picture
plt.plot(id, w1, label='Trend of w1')
plt.xlabel('id');
plt.ylabel('w1');
plt.legend();
plt.show();

plt.plot(id, w2, label='Trend of w2')
plt.xlabel('id');
plt.ylabel('w2');
plt.ylim([0.0, 0.5]);
plt.legend();
plt.show();
```



Yes.

w_1 approaches 2.78. The converging speed is fast at first, then becomes really slow before w_1 is stabilized.

w_2 approaches 0. The converging speed is fast at first, then becomes really slow before w_2 is stabilized.

In []: