

Problem 1

In [4]:

```
import autograd.numpy as ag_np

# Use helper packages
from AbstractBaseCollabFilterSGD import AbstractBaseCollabFilterSGD
from utils import load_dataset
import numpy as np
# Some packages you might need (uncomment as necessary)
import pandas as pd
import matplotlib.pyplot as plt
import os
```

In [4]:

```
from CollabFilterMeanOnly import CollabFilterMeanOnly
train_tuple, valid_tuple, test_tuple, n_users, n_items = load_dataset()
n_epochs=50
model = CollabFilterMeanOnly(n_epochs=n_epochs)
model.init_parameter_dict(n_users, n_items, train_tuple)
model.fit(train_tuple, valid_tuple)
```


poch	0.000	loss_total	13.82200	train_MAE	3.53239	valid
_MAE	3.50620	grad_wrt_mu	7.10400			
poch	0.013	loss_total	9.11851	train_MAE	2.82199	valid
_MAE	2.79580	grad_wrt_mu	5.63320			
poch	0.025	loss_total	6.60551	train_MAE	2.29207	valid
_MAE	2.26637	grad_wrt_mu	4.63456			
poch	0.100	loss_total	1.68633	train_MAE	1.05029	valid
_MAE	1.03490	grad_wrt_mu	1.24038			
poch	0.200	loss_total	1.26837	train_MAE	0.95532	valid
_MAE	0.95363	grad_wrt_mu	0.10245			
poch	0.313	loss_total	1.28873	train_MAE	0.94556	valid
_MAE	0.94652	grad_wrt_mu	0.04033			
poch	0.400	loss_total	1.32975	train_MAE	0.94431	valid
_MAE	0.94560	grad_wrt_mu	0.02424			
poch	0.500	loss_total	1.21539	train_MAE	0.94557	valid
_MAE	0.94652	grad_wrt_mu	0.00853			
poch	0.613	loss_total	1.27747	train_MAE	0.94446	valid
_MAE	0.94571	grad_wrt_mu	0.01354			
poch	0.713	loss_total	1.24705	train_MAE	0.94475	valid

In [17]:

```
plt.figure
plt.plot(model.trace_epoch, model.trace_mae_train, label='Tr')
plt.plot(model.trace_epoch, model.trace_mae_valid, label='Va')
plt.title('MAE Behavior')
plt.xlabel('Epoch')
plt.ylabel('MAE')
plt.legend()
```

Out[17]:

<matplotlib.legend.Legend at 0x1a88bf8cf98>

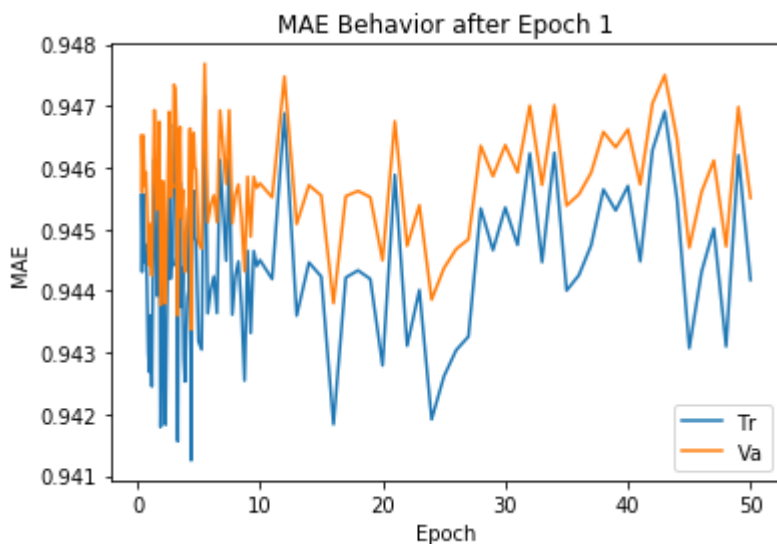


In [18]:

```
plt.figure
plt.xlabel('Epoch')
plt.ylabel('MAE')
plt.title('MAE Behavior after Epoch 1')
plt.plot(model.trace_epoch[5:], model.trace_mae_train[5:], label='Tr')
plt.plot(model.trace_epoch[5:], model.trace_mae_valid[5:], label='Va')
plt.legend()
```

Out[18]:

<matplotlib.legend.Legend at 0x1a88bff7d68>



Questions:

1b No, we are using mean as prediction. Forcing μ to be small by adding penalty, doesn't help the prediction at all. In fact, it's driving μ away from the optimal result.

What we need to consider is whether a parameter implicates complexity of a model so we want to penalize it. Statistical parameters like mean and var are not a measure of complexity and we don't want to penalize it.

1c The actual average is 3.53239. The SGD prediction is 3.53671, which agrees with the actual ave.

In [24]:

```
print(np.mean(train_tuple[2]), model.param_dict["mu"][0])
```

3.5323907390739073 3.536711325075617

In []:

Problem2

In [10]:

```
from CollabFilterOneScalarPerItem import CollabFilterOneScalarPerItem
```

In [11]:

```
train_tuple, valid_tuple, test_tuple, n_users, n_items = load_dataset()
model2 = CollabFilterOneScalarPerItem(
    n_epochs=250, step_size=0.5)
model2.init_parameter_dict(n_users, n_items, train_tuple)
model2.fit(train_tuple, valid_tuple)
```

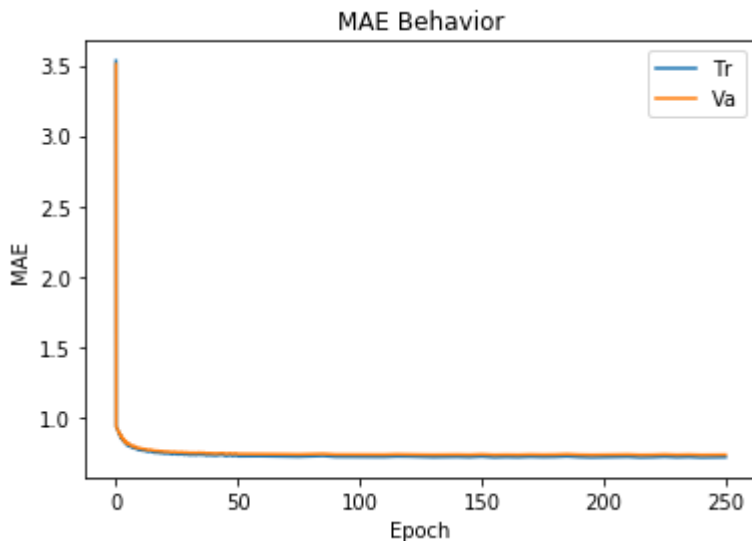
poch	0.000		loss_total	13.82200		train_MAE	3.53239		valid
_MAE	3.50620		grad_wrt_mu	7.10400		grad_wrt_b_per_user	0.007		
3		grad_wrt_c_per_item	0.00423						
poch	0.013		loss_total	1.18426		train_MAE	0.93962		valid
_MAE	0.94187		grad_wrt_mu	0.07603		grad_wrt_b_per_user	0.001		
8		grad_wrt_c_per_item	0.00082						
poch	0.025		loss_total	1.23511		train_MAE	0.94317		valid
_MAE	0.94429		grad_wrt_mu	0.12661		grad_wrt_b_per_user	0.001		
2		grad_wrt_c_per_item	0.00084						
poch	0.100		loss_total	1.28834		train_MAE	0.93797		valid
_MAE	0.93971		grad_wrt_mu	0.01212		grad_wrt_b_per_user	0.001		
4		grad_wrt_c_per_item	0.00088						
poch	0.200		loss_total	1.24465		train_MAE	0.93125		valid
_MAE	0.93359		grad_wrt_mu	0.14125		grad_wrt_b_per_user	0.001		
6		grad_wrt_c_per_item	0.00086						
poch	0.313		loss_total	1.24821		train_MAE	0.92916		valid
_MAE	0.93086		grad_wrt_mu	0.02762		grad_wrt_b_per_user	0.001		
6		grad_wrt_c_per_item	0.00088						
poch	0.400		loss_total	1.26587		train_MAE	0.92292		valid

In [12]:

```
plt.figure
plt.plot(model2.trace_epoch, model2.trace_mae_train, label='Tr')
plt.plot(model2.trace_epoch, model2.trace_mae_valid, label='Va')
plt.title('MAE Behavior')
plt.xlabel('Epoch')
plt.ylabel('MAE')
plt.legend()
```

Out[12]:

<matplotlib.legend.Legend at 0x1d12e3fa7b8>

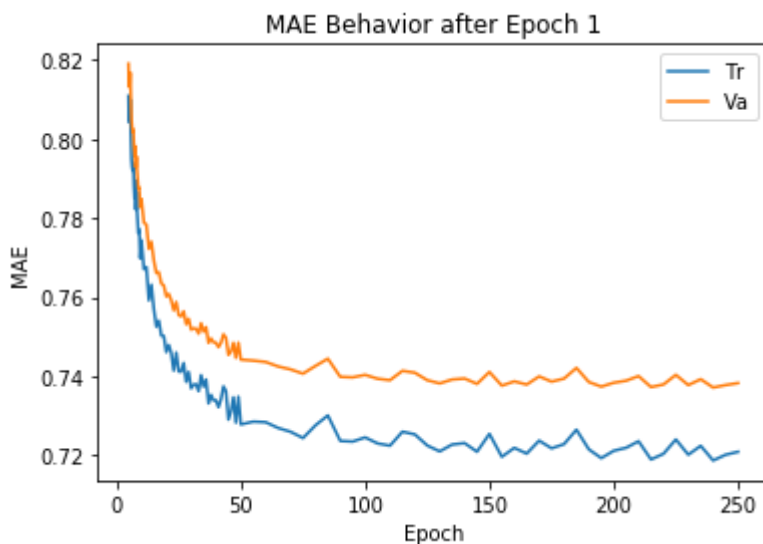


In [13]:

```
plt.figure
plt.xlabel('Epoch')
plt.ylabel('MAE')
plt.title('MAE Behavior after Epoch 1')
plt.plot(model2.trace_epoch[50:], model2.trace_mae_train[50:], label='Tr')
plt.plot(model2.trace_epoch[50:], model2.trace_mae_valid[50:], label='Va')
plt.legend()
```

Out[13]:

<matplotlib.legend.Legend at 0x1d12dda95c0>



Questions:

2a The MAE performance on the Validation set improved by over 20% (from 0.945 to 0.741).

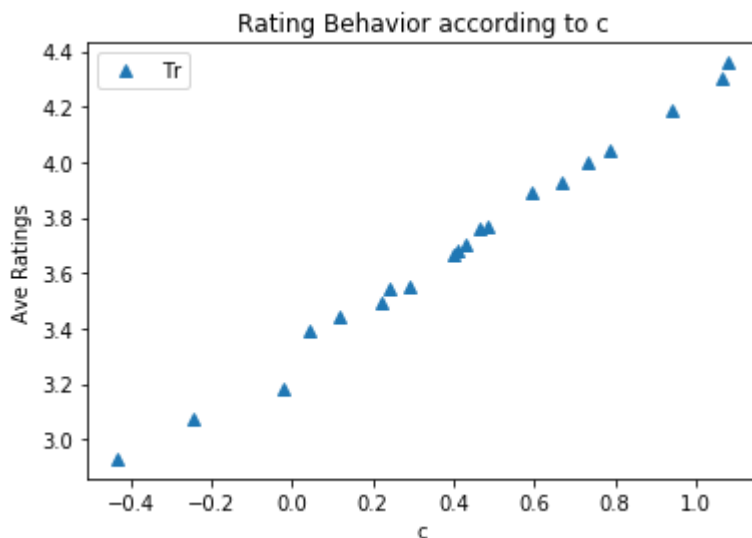
2b From the Picture Below, rating seems to have a linear dependence on trained c . When c is larger, the average rating of a movie seems higher. When c is small (very negative), the average rating is lower.

In [14]:

```
#2b
data_path='C:/Users/xush4/Documents/comp135-19s-assignments-master/project3/data_movie_lens
sel_df = pd.read_csv(os.path.join(data_path, "select_movies.csv"))
aveT=[]; c_sel=[]
for i in sel_df["item_id"]:
    idx=train_tuple[1]
    aveTid=np.mean(train_tuple[2][np.where(idx==i)])
    aveT.append(aveTid);
    c_sel.append(model2.param_dict["c_per_item"][i])
    ##print(i, model2.param_dict["c_per_item"][i], aveid)
plt.plot(c_sel, aveT, linestyle='', marker='^', label="Tr")
plt.xlabel('c')
plt.ylabel('Ave Ratings')
plt.title('Rating Behavior according to c')
plt.legend()
```

Out[14]:

<matplotlib.legend.Legend at 0x1d12c10c9e8>



In []:

Problem3

In [6]:

```
from CollabFilterOneVectorPerItem import CollabFilterOneVectorPerItem as CFV
```

In [4]:

```
K=[0,2,10,50]
model3={};
for i in range(4):
    train_tuple, valid_tuple, test_tuple, n_users, n_items = load_dataset()
    model3[i] = CFV(n_epochs=250, step_size=0.5, n_factors=K[i])
    model3[i].init_parameter_dict(n_users, n_items, train_tuple)
    model3[i].fit(train_tuple, valid_tuple)
```

```
:\\ProgramData\\Anaconda3\\lib\\site-packages\\numpy\\core\\fromnumeric.py:3118:
untimeWarning: Mean of empty slice.
    out=out, **kwargs)
:\\ProgramData\\Anaconda3\\lib\\site-packages\\numpy\\core\\_methods.py:85: Runt
meWarning: invalid value encountered in double_scalars
    ret = ret.dtype.type(ret / rcount)
```

In [16]:

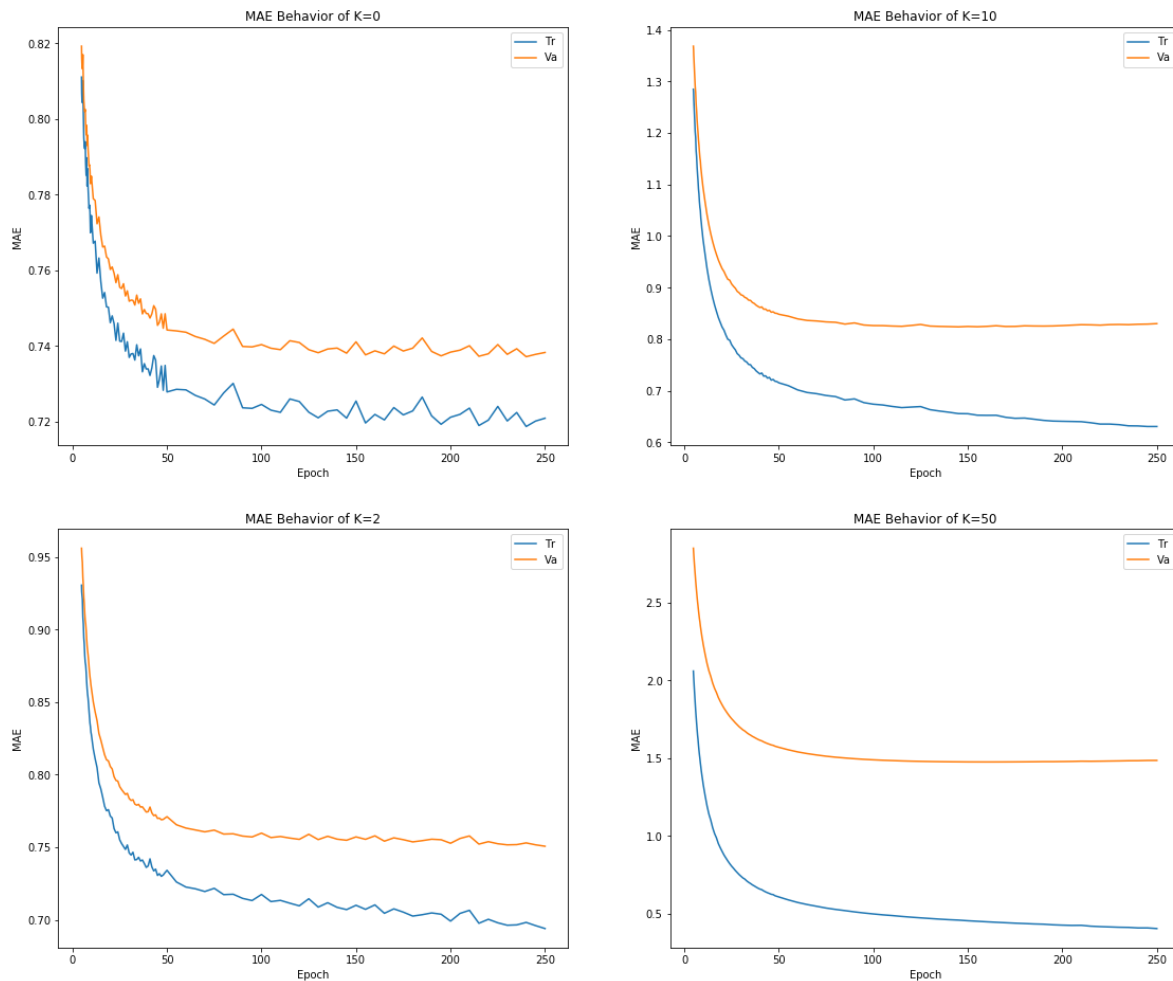
```
fig3a, axes = plt.subplots(nrows=2, ncols=2, figsize=(19,16))
for i in range(2):
    for j in range(2):
        axes[i,j].set_xlabel('Epoch')
        axes[i,j].set_ylabel('MAE')
        axes[i,j].set_title('MAE Behavior of K=' + str(int(K[i+2*j])))
        axes[i,j].plot(model3[i+2*j].trace_epoch[50:], model3[i+2*j].trace_mae_train[50:],
        axes[i,j].plot(model3[i+2*j].trace_epoch[50:], model3[i+2*j].trace_mae_valid[50:],
        axes[i,j].legend()

fig3a.suptitle('MAE behavior,'+ 'Alpha='+str(0))
```

Out[16]:

Text(0.5,0.98,'MAE behavior,Alpha=0')

MAE behavior,Alpha=0



In [7]:

```
### 3b
K=[0,2,10,50]
a=np.logspace(-10,2)
tr_error=[];
va_error=[];
for i in range(4):
    te=[]; ve=[];
    for j in range(int(a.size)):
        train_tuple, valid_tuple, test_tuple, n_users, n_items = load_dataset()
        model3b = CFV(n_epochs=250, step_size=0.5, n_factors=K[i], alpha=a[j])
        model3b.init_parameter_dict(n_users, n_items, train_tuple)
        model3b.fit(train_tuple, valid_tuple)
        te.append(model3b.trace_mae_train[-1])
        ve.append(model3b.trace_mae_valid[-1])
    tr_error.append(te)
    va_error.append(ve)
```

```
:\\ProgramData\\Anaconda3\\lib\\site-packages\\numpy\\core\\fromnumeric.py:3118:
untimeWarning: Mean of empty slice.
    out=out, **kwargs)
:\\ProgramData\\Anaconda3\\lib\\site-packages\\numpy\\core\\_methods.py:85: Runt
meWarning: invalid value encountered in double_scalars
    ret = ret.dtype.type(ret / rcount)
```


In [8]:

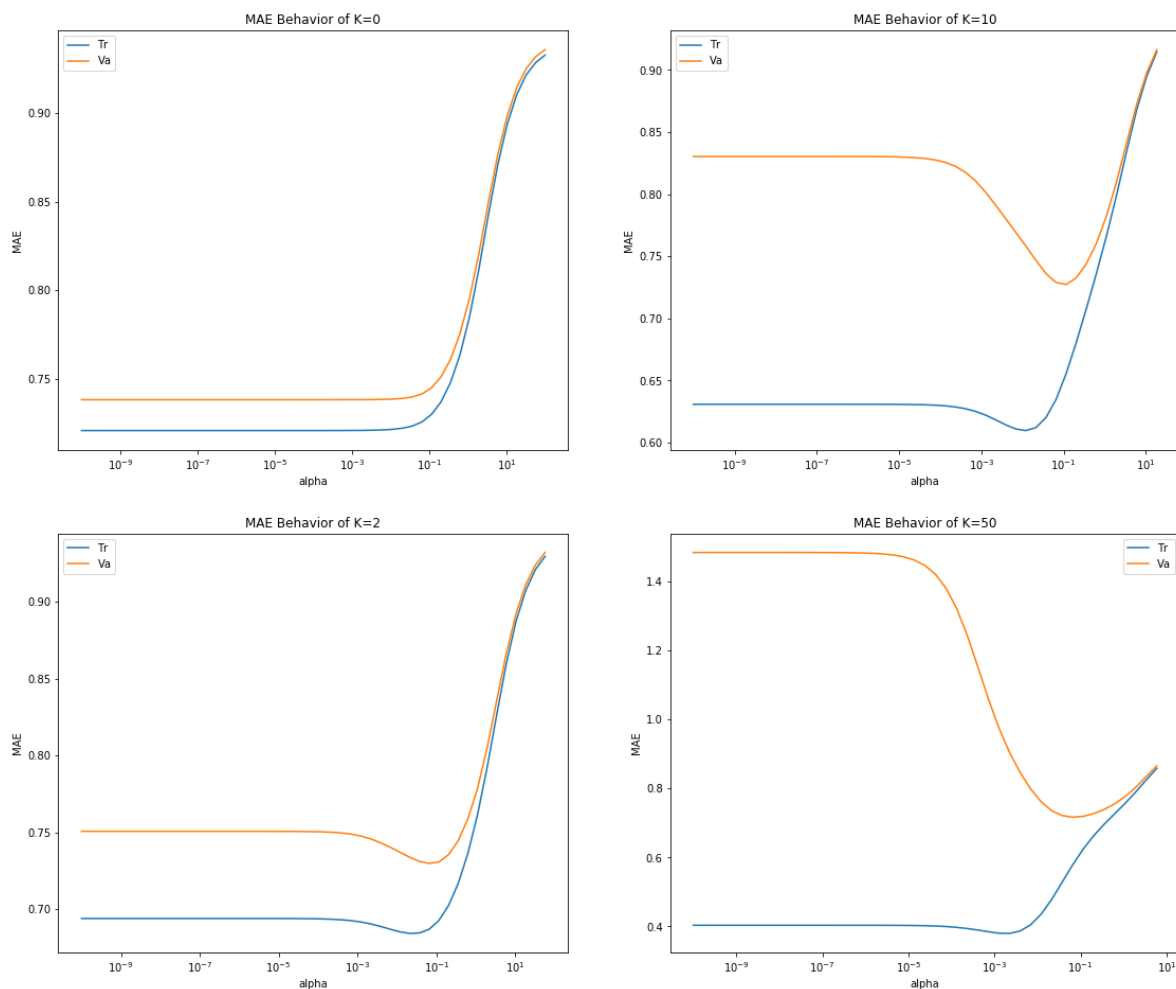
```
fig3a, axes = plt.subplots(nrows=2, ncols=2, figsize=(19,16))
for i in range(2):
    for j in range(2):
        axes[i,j].set_xlabel('alpha')
        axes[i,j].set_ylabel('MAE')
        axes[i,j].set_title('MAE Behavior of K=' + str(int(K[i+2*j])))
        axes[i,j].plot(a, tr_error[i+2*j], label='Tr')
        axes[i,j].plot(a, va_error[i+2*j], label='Va')
        axes[i,j].legend()
        axes[i,j].set_xscale('log')

fig3a.suptitle('MAE behavior, '+ 'Alpha=np.logspace(-10,5)')
```

Out[8]:

Text(0.5,0.98,'MAE behavior,Alpha=np.logspace(-10,5)')

MAE behavior,Alpha=np.logspace(-10,5)



3c We can use a small K or shrink the size of each batch.

3d K=0 and K=2 are performing a little better than the model in question 1 and 2.

I would recommend use 0 or 2 factors.

There's no need to look into more than 50 factors because it's not increasing the performance(probably overfitting).

In [7]:

#2 Factors:

```

train_tuple, valid_tuple, test_tuple, n_users, n_items = load_dataset()
model3e = CFV(n_epochs=250, step_size=0.5, n_factors=2, alpha=0.1)
model3e.init_parameter_dict(n_users, n_items, train_tuple)
model3e.fit(train_tuple, valid_tuple)

```

```

poch      0.000 | loss_total      6.21473 | train_MAE      1.50279 | valid
_MAE      1.49613 | grad_wrt_mu      0.92071 | grad_wrt_b_per_user      0.002
7 | grad_wrt_c_per_item      0.00149 | grad_wrt_U      0.00348 | grad_wrt_V
.00234
poch      0.013 | loss_total      6.17911 | train_MAE      1.41942 | valid
_MAE      1.42198 | grad_wrt_mu      0.12156 | grad_wrt_b_per_user      0.002
9 | grad_wrt_c_per_item      0.00141 | grad_wrt_U      0.00342 | grad_wrt_V
.00239
poch      0.025 | loss_total      6.03396 | train_MAE      1.41146 | valid
_MAE      1.41516 | grad_wrt_mu      0.03582 | grad_wrt_b_per_user      0.002
9 | grad_wrt_c_per_item      0.00138 | grad_wrt_U      0.00333 | grad_wrt_V
.00230
poch      0.100 | loss_total      5.76014 | train_MAE      1.39247 | valid
_MAE      1.39530 | grad_wrt_mu      0.19923 | grad_wrt_b_per_user      0.002
5 | grad_wrt_c_per_item      0.00135 | grad_wrt_U      0.00326 | grad_wrt_V
.00225
poch      0.200 | loss_total      5.32516 | train_MAE      1.35524 | valid
_MAE      1.36099 | grad_wrt_mu      0.01691 | grad_wrt_b_per_user      0.002
0 | grad_wrt_c_per_item      0.00132 | grad_wrt_U      0.00293 | grad_wrt_V

```

In [28]:

```

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

data_path='C:/Users/xush4/Documents/comp135-19s-assignments-master/project3/data_movie_lens
sel_df = pd.read_csv(os.path.join(data_path, "select_movies.csv"))
aveT=[]; index=[]
for i in sel_df["item_id"]:
    idx=train_tuple[1]
    aveTid=np.mean(train_tuple[2][np.where(idx==i)])
    aveT.append(aveTid);
    index.append(i)

V=model3e.param_dict['V'][index]
#print(V[:,0].size, V[:,1].size, aveT)
ax.scatter(V[:,0], V[:,1], aveT, zdir='z')


ax.set_xlabel('V[1]')
ax.set_ylabel('V[2]')
ax.set_zlabel('Score')
print(V[:,0], V[:,1], aveT)

```

```

0.08239941 -0.32913701 -0.0372406 0.22327492 -0.11569277 -0.32949667
-0.51060424 -0.50746085 -0.79073236 -0.592902 -0.65224791 -0.64892778
-0.80690433 -0.80711085 -0.6925586 0.54269781 0.06421546 -0.53089196
-0.50387622 -0.48682363 [ 0.3561754 0.53450276 -0.15376184 -0.04260292
.77782893 -0.15445973
-0.19469097 0.1935403 0.25841157 0.19861086 0.13559631 0.2752468
0.02343947 0.09642642 -0.13832645 -0.27475537 -0.48355823 0.1091072
-0.07069955 -0.05248082 [3.888888888888889, 3.76271186440678, 3.7007299270
72993, 4.045, 3.67816091954023, 4.359574468085106, 4.191335740072202, 4.002
50980392157, 3.6666666666666665, 2.9302325581395348, 4.306990881458966, 3.9
8030303030303, 3.546153846153846, 3.4939759036144578, 3.3943661971830985,
.5535714285714284, 3.7710843373493974, 3.180722891566265, 3.441624365482233
, 3.076923076923077]

```



The trend is not that obvious, it seems to me items with smaller V[1] and bigger V[2] are more controversial(either very high scores or very low scores.).

Problem 4

In [160]:

```
# Reference https://surprise.readthedocs.io/en/stable/matrix\_factorization.html#surprise.pr  
from surprise.prediction_algorithms.matrix_factorization import SVD
```

In [149]:

```
import pandas as pd  
import numpy as np  
  
from surprise import SVD  
from surprise import Dataset, Reader, accuracy  
from surprise.model_selection import cross_validate, KFold  
  
reader = Reader(  
    line_format='user item rating', sep=',',  
    rating_scale=(1, 5), skip_lines=1)  
  
## Load the training set into surprise's custom dataset object  
train_df = pd.read_csv('data_movie_lens_100k/ratings_train.csv')  
train_set = Dataset.load_from_file('data_movie_lens_100k/ratings_train.csv', reader=reader)  
  
## Load the test set into surprise's custom dataset object  
## (Need to use intermediate pandas DataFrame because the true ratings are missing)  
test_df = pd.read_csv('data_movie_lens_100k/ratings_test_masked.csv')  
test_set = Dataset.load_from_df(test_df, reader=reader)  
test_set = test_set.build_full_trainset().build_testset()  
print(type(test_df['user_id'][0]))  
  
<class 'numpy.int64'>
```

In [48]:

```
numF=5
kf = KFold(n_splits=numF)
K=[0,2,5,10]
a=np.logspace(-3,1,17)
tr_error=[];
va_error=[];

for i in range(4):
    te=[]; ve=[];
    for j in range(int(a.size)):
        sumTe=0; sumVa=0;
        for trainset, validset in kf.split(train_set):
            model4 = SVD(n_epochs=250, n_factors=K[i], lr_all=a[j])
            # train and test algorithm.
            model4.fit(trainset)
            pre1 = model4.test(trainset.build_testset())
            pre2 = model4.test(validset)
            # Compute and print Root Mean Squared Error
            sumTe=sumTe+accuracy.mae(pre1, verbose=True)
            sumVa=sumVa+accuracy.mae(pre2, verbose=True)
        print(sumTe/numF)
        te.append(sumTe/numF);
        ve.append(sumVa/numF);
    tr_error.append(te)
    va_error.append(ve)
```

```
MAE: 0.5863
MAE: 0.7427
MAE: 0.5889
MAE: 0.7405
MAE: 0.5840
MAE: 0.7499
MAE: 0.5906
MAE: 0.7426
0.5869958185262447
MAE: 0.5368
MAE: 0.7519
MAE: 0.5346
MAE: 0.7572
MAE: 0.5365
MAE: 0.7670
MAE: 0.5365
MAE: 0.7524
MAE: 0.5381
MAE: 0.7549
```

In [49]:


```
#print(tr_error)
fig3a, axes = plt.subplots(nrows=2, ncols=2, figsize=(19,16))
for i in range(2):
    for j in range(2):
        axes[i,j].set_xlabel('alpha')
        axes[i,j].set_ylabel('MAE')
        axes[i,j].set_title('MAE Behavior of K=' + str(int(K[i+2*j])))
        axes[i,j].plot(a, tr_error[i+2*j], label='Tr')
        axes[i,j].plot(a, va_error[i+2*j], label='Va')
        axes[i,j].legend()
        axes[i,j].set_xscale('log')

fig3a.suptitle('MAE behavior, '+ 'Alpha=np.logspace(-3,2)')
```

Out[49]:

Text(0.5,0.98,'MAE behavior,Alpha=np.logspace(-3,2)')

MAE behavior,Alpha=np.logspace(-3,2)



The trend on $\alpha \leq 0.1$ is pretty identical. The trend over $\alpha > 0.1$ is a differnt. It's probably because the step size we use is larger so the method can not converge near enough to the optimal solution.

In [202]:

```
reader = Reader(
    line_format='user item rating', sep=',',
    rating_scale=(1, 5), skip_lines=1)

## Load the training set into surprise's custom dataset object
## (Need to use intermediate pandas DataFrame here because that's what needed on test set)
train_df = pd.read_csv('data_movie_lens_100k/ratings_train.csv')
train_set = Dataset.load_from_df(train_df, reader=reader)
train_set = train_set.build_full_trainset()

## Load the test set into surprise's custom dataset object
## (Need to use intermediate pandas DataFrame because the true ratings are missing)
test_df = pd.read_csv('data_movie_lens_100k/ratings_test_masked.csv')
test_set = Dataset.load_from_df(test_df, reader=reader)
test_set = test_set.build_full_trainset().build_testset()

# Use the SVD algorithm
    ## Fit model to training set
model = SVD(n_factors=2)
model.fit(train_set)

## Measure predictions on train set
print("Making predictions on training set (showing first 10):")
tr_pred = model.test(train_set.build_testset())
tr_mae = accuracy.mae(tr_pred)
tr_predicted_ratings_N = np.asarray([p.est for p in tr_pred], dtype=np.float64)
print(tr_predicted_ratings_N[:10])

## Measure predictions on test set
print("Making predictions on test set (showing first 10):")
te_pred = model.test(test_set)
te_mae = accuracy.mae(te_pred) # should be NaN because no real labels on testset
te_predicted_ratings_N = np.asarray([p.est for p in te_pred], dtype=np.float64)
print(te_predicted_ratings_N[:10])

print("n_factors %6d  tr_MAE %7.3f  test_MAE %7.3f" % (n_factors, tr_mae, te_mae))

print("Making test set predictions in the original order")
for row in test_df.values[:10]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    print("user %4d  item %4d  predicted rating % 8.3f" % (userid, itemid, rhat.est))
tep0=[]
for row in test_df.values[:]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    #print("user %4d  item %4d  predicted rating % 8.3f" % (userid, itemid, rhat.est))
    tep0.append(rhat.est)
np.savetxt('predicted_ratings_test4.txt', np.asarray(tep0))
```

Making predictions on training set (showing first 10):

MAE: 0.7184

[2.69729623 3.61916657 2.80442954 3.72018483 2.77842302 3.47178856
3.28690477 3.43117354 3.70169473 2.63116753]

Making predictions on test set (showing first 10):

MAE: nan

[4.20307785 3.23106528 2.51561078 3.45627649 3.97710067 3.86847339

```
2.98476368 3.7159181 3.16108109 3.64082798]
n_factors      2  tr_MAE    0.718  test_MAE      nan
Making test set predictions in the original order
user  503  item  204  predicted rating    4.203
user  795  item  185  predicted rating    3.908
user   42  item  403  predicted rating    3.751
user  327  item  740  predicted rating    3.450
user  285  item   98  predicted rating    3.989
user  279  item   11  predicted rating    4.483
user  496  item  588  predicted rating    2.373
user  499  item  266  predicted rating    3.572
user  357  item  126  predicted rating    4.146
user  932  item  182  predicted rating    3.213
```

Problem5

In [50]:

```
#Kmeans:
```

In [196]:

```
numF=5
kf = KFold(n_splits=numF)
train_df = pd.read_csv('data_movie_lens_100k/ratings_train.csv')
train_set = Dataset.load_from_file('data_movie_lens_100k/ratings_train.csv', reader=reader)

total_sample=train_df.shape[0]
K_nbh=((1-1/numF)*total_sample)**np.linspace(0.25, 0.75)
K_nbh=[1,5,10,25,50, int(np.sqrt((1-1/numF)*total_sample))]
```

In [90]:

```
print(K_nbh)
```

```
[1, 5, 10, 25, 50, 268]
```


In [91]:

```
##MSD with mean
from surprise.prediction_algorithms.knns import KNNWithMeans as KNNM
te=[]; ve=[];
for k in K_nbh:
    sumTe=0; sumVa=0;
    for trainset, validset in kf.split(train_set):
        model5=KNNM(k=int(k))
        # train and test algorithm.
        model5.fit(trainset)
        pre1 = model5.test(trainset.build_testset())
        pre2 = model5.test(validset)
        sumTe=sumTe+accuracy.mae(pre1, verbose=True)
        sumVa=sumVa+accuracy.mae(pre2, verbose=True)
    te.append(sumTe/numF);
    ve.append(sumVa/numF);
```

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0522

MAE: 0.9743

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0491

MAE: 0.9796

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0504

MAE: 0.9749

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0492

MAE: 0.9698

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0512

In [197]:

```
##MSD with mean
from surprise.prediction_algorithms.knns import KNNWithMeans as KNNM
tea=[]; vea=[];
for k in K_nbh:
    sumTe=0; sumVa=0;
    for trainset, validset in kf.split(train_set):
        model5=KNNM(k=int(k),sim_options={'name': 'pearson'})
        # train and test algorithm.
        model5.fit(trainset)
        pre1 = model5.test(trainset.build_testset())
        pre2 = model5.test(validset)
        sumTe=sumTe+accuracy.mae(pre1, verbose=True)
        sumVa=sumVa+accuracy.mae(pre2, verbose=True)
    tea.append(sumTe/numF);
    vea.append(sumVa/numF);
```

Computing the pearson similarity matrix...

Done computing similarity matrix.

MAE: 0.1289

MAE: 0.9988

Computing the pearson similarity matrix...

Done computing similarity matrix.

MAE: 0.1305

MAE: 1.0057

Computing the pearson similarity matrix...

Done computing similarity matrix.

MAE: 0.1313

MAE: 1.0045

Computing the pearson similarity matrix...

Done computing similarity matrix.

MAE: 0.1326

MAE: 1.0004

Computing the pearson similarity matrix...

Done computing similarity matrix.

MAE: 0.1307

In [92]:

```
##MSD
from surprise.prediction_algorithms.knns import KNNBasic as KNN
teb=[]; veb=[];
for k in K_nbh:
    sumTe=0; sumVa=0;
    for trainset, validset in kf.split(train_set):
        model5=KNN(k=int(k))
        # train and test algorithm.
        model5.fit(trainset)
        pre1 = model5.test(trainset.build_testset())
        pre2 = model5.test(validset)
        sumTe=sumTe+accuracy.mae(pre1, verbose=True)
        sumVa=sumVa+accuracy.mae(pre2, verbose=True)
    teb.append(sumTe/numF);
    veb.append(sumVa/numF);
```

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0000

MAE: 0.9784

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0000

MAE: 0.9739

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0000

MAE: 0.9722

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0000

MAE: 0.9740

Computing the msd similarity matrix...

Done computing similarity matrix.

MAE: 0.0000

In [93]:

```

## Pearson
from surprise.prediction_algorithms.knns import KNNBasic as KNN
tec=[]; vec=[];

for k in K_nbh:
    sumTe=0; sumVa=0;
    for trainset, validset in kf.split(train_set):
        model5=KNN(k=int(k),sim_options={'name': 'pearson'})
        # train and test algorithm.
        model5.fit(trainset)
        pre1 = model5.test(trainset.build_testset())
        pre2 = model5.test(validset)
        sumTe=sumTe+accuracy.mae(pre1, verbose=True)
        sumVa=sumVa+accuracy.mae(pre2, verbose=True)
    tec.append(sumTe/numF);
    vec.append(sumVa/numF);

```

```

Computing the pearson similarity matrix...
Done computing similarity matrix.
MAE: 0.1378
MAE: 1.0490
Computing the pearson similarity matrix...
Done computing similarity matrix.
MAE: 0.1337
MAE: 1.0582
Computing the pearson similarity matrix...
Done computing similarity matrix.
MAE: 0.1337
MAE: 1.0559
Computing the pearson similarity matrix...
Done computing similarity matrix.
MAE: 0.1344
MAE: 1.0568
Computing the pearson similarity matrix...
Done computing similarity matrix.
MAE: 0.1350

```

Summary of 5:

I use surprise KNN for problem 5. I look into the traditional KNN, KNN-Means(KNNM, subtract means), KNN using pearson correlation(KNNP) and KNN-Means with pearson correlation(KNNMP). I looked into how k affects the error of the methods. It seems choosing K=20 should be a good strategy in all of the cases. We observe the test error behavior as the following:


KNN: 0.7653 KNNM: 0.7427 KNNP: 0.7953 KNNMP:0.7367 This shows when using KNN here, subtract mean and use pearson correlation is better for prediction.

In [94]:

```
plt.figure
plt.xlabel('K')
plt.ylabel('MAE')
plt.title('KNN-means MAE Behavior')
plt.plot(K_nbh, te, label='Tr')
plt.plot(K_nbh, ve, label='Va')
plt.xscale('log')
plt.legend()
```

Out[94]:

<matplotlib.legend.Legend at 0x20e89fc7748>

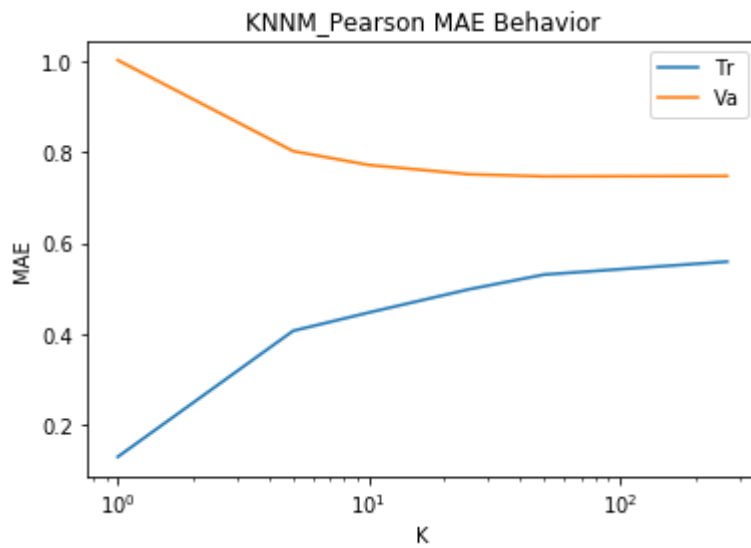


In [198]:

```
plt.figure
plt.xlabel('K')
plt.ylabel('MAE')
plt.title('KNNM_Pearson MAE Behavior')
plt.plot(K_nbh, tea, label='Tr')
plt.plot(K_nbh, vea, label='Va')
plt.xscale('log')
plt.legend()
```

Out[198]:

<matplotlib.legend.Legend at 0x20e8a34f6d8>

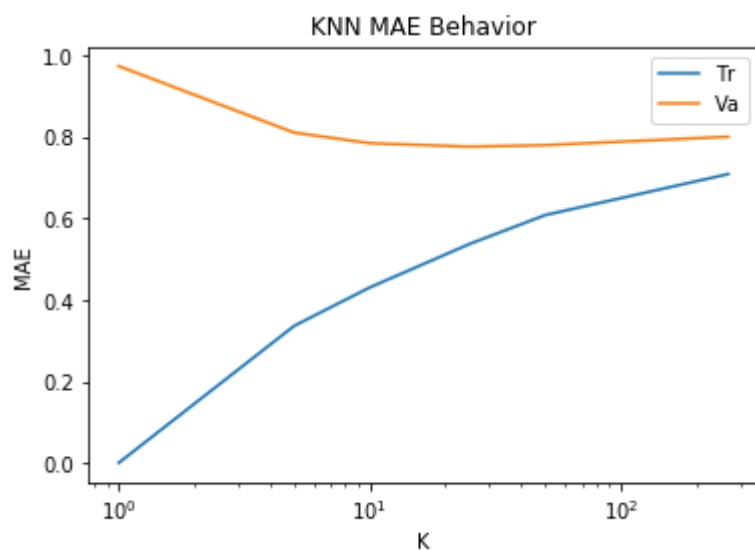


In [95]:

```
plt.figure
plt.xlabel('K')
plt.ylabel('MAE')
plt.title('KNN MAE Behavior')
plt.plot(K_nbh, teb, label='Tr')
plt.plot(K_nbh, veb, label='Va')
plt.xscale('log')
plt.legend()
```

Out[95]:

<matplotlib.legend.Legend at 0x20e8a36f780>




In [96]:

```
plt.figure
plt.xlabel('K')
plt.ylabel('MAE')
plt.title('KNN-pearson MAE Behavior')
plt.plot(K_nbh, tec, label='Tr')
plt.plot(K_nbh, vec, label='Va')
plt.xscale('log')
plt.legend()
```

Out[96]:

<matplotlib.legend.Legend at 0x20e880df5c0>



In [200]:

```
### Comming back to get result
train_set = train_set.build_full_trainset()
```


In [189]:

```

model=KNNM(K=20)
model.fit(train_set)

## Measure predictions on train set
print("Making predictions on training set (showing first 10):")
tr_pred = model.test(train_set.build_testset())
tr_mae = accuracy.mae(tr_pred)
tr_predicted_ratings_N = np.asarray([p.est for p in tr_pred], dtype=np.float64)
print(tr_predicted_ratings_N[:10])

## Measure predictions on test set
print("Making predictions on test set (showing first 10):")
te_pred = model.test(test_set)
#te_mae = accuracy.mae(te_pred) # should be NaN because no real labels on testset
te_predicted_ratings_N1 = np.asarray([p.est for p in te_pred], dtype=np.float64)
#print(te_predicted_ratings_N[:10])

print("n_factors %d tr_MAE %7.3f test_MAE %7.3f" % (n_factors, tr_mae, te_mae))

#print("Making test set predictions in the original order")
tep1=[]
for row in test_df.values[:]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    #print("user %4d item %4d predicted rating % 8.3f" % (userid, itemid, rhat.est))
    tep1.append(rhat.est)
np.savetxt('predicted_ratings_test5M.txt', np.asarray(tep1))

```

Computing the msd similarity matrix...

Done computing similarity matrix.

Making predictions on training set (showing first 10):

MAE: 0.5981

```

[2.62806394 3.51219776 2.90599123 3.98393553 2.89749724 2.91080894
 3.57222821 3.2620783 3.64707388 2.3300006 ]

```

Making predictions on test set (showing first 10):

MAE: nan

n_factors	2	tr_MAE	0.598	test_MAE	nan
-----------	---	--------	-------	----------	-----

In [190]:

```

#train_set = train_set.build_full_trainset()
model=KNN(K=20)
model.fit(train_set)

## Measure predictions on train set
print("Making predictions on training set (showing first 10):")
tr_pred = model.test(train_set.build_testset())
tr_mae = accuracy.mae(tr_pred)
tr_predicted_ratings_N = np.asarray([p.est for p in tr_pred], dtype=np.float64)
print(tr_predicted_ratings_N[:10])

## Measure predictions on test set
print("Making predictions on test set (showing first 10):")
te_pred = model.test(test_set)
#te_mae = accuracy.mae(te_pred) # should be NaN because no real labels on testset
te_predicted_ratings_N2 = np.asarray([p.est for p in te_pred], dtype=np.float64)
#print(te_predicted_ratings_N2[:10])

print("n_factors %d  tr_MAE %7.3f  test_MAE %7.3f" % (n_factors, tr_mae, te_mae))

#print("Making test set predictions in the original order")
tep2=[]
for row in test_df.values[:]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    #print("user %4d item %4d predicted rating % 8.3f" % (userid, itemid, rhat.est))
    tep2.append(rhat.est)
np.savetxt('predicted_ratings_test5.txt', np.asarray(tep2))

```

Computing the msd similarity matrix...

Done computing similarity matrix.

Making predictions on training set (showing first 10):

MAE: 0.6015

```

[2.47094827 3.91616195 3.00816135 4.30859069 3.07263422 3.1585108
 3.73761608 3.51404379 3.84158116 2.40704005]

```

Making predictions on test set (showing first 10):

MAE: nan

```

n_factors      2  tr_MAE    0.602  test_MAE      nan

```

In [191]:

```

model=KNN(K=20, sim_options={'name': 'pearson'})
model.fit(train_set)

## Measure predictions on train set
print("Making predictions on training set (showing first 10):")
tr_pred = model.test(train_set.build_testset())
tr_mae = accuracy.mae(tr_pred)
tr_predicted_ratings_N = np.asarray([p.est for p in tr_pred], dtype=np.float64)
print(tr_predicted_ratings_N[:10])

## Measure predictions on test set
print("Making predictions on test set (showing first 10):")
te_pred = model.test(test_set)
#te_mae = accuracy.mae(te_pred) # should be NaN because no real labels on testset
te_predicted_ratings_N = np.asarray([p.est for p in te_pred], dtype=np.float64)
print(te_predicted_ratings_N[:10])

print("n_factors %d  tr_MAE %7.3f  test_MAE %7.3f" % (n_factors, tr_mae, te_mae))

#print("Making test set predictions in the original order")
tep3=[]
for row in test_df.values[:]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    #print("user %4d  item %4d  predicted rating % 8.3f" % (userid, itemid, rhat.est))
    tep3.append(rhat.est)
np.savetxt('predicted_ratings_test5p.txt', np.asarray(tep3))

```

Computing the pearson similarity matrix...

Done computing similarity matrix.

Making predictions on training set (showing first 10):

MAE: 0.6067

```
[2.46275582 4.12447317 3.10790543 4.44937415 3.13762163 3.51988639
 3.83512304 3.43602947 3.96805273 2.54442012]
```

Making predictions on test set (showing first 10):

MAE: nan

```
[3.81880269 2.38702228 2.33748343 3.77718842 3.93141071 3.53435695
 3.0212087 3.80168179 3.0288308 3.69002411]
```

```
n_factors      2  tr_MAE    0.607  test_MAE      nan
```

In [201]:

```

model=KNNM(K=20, sim_options={'name': 'pearson'})
model.fit(train_set)

## Measure predictions on train set
print("Making predictions on training set (showing first 10):")
tr_pred = model.test(train_set.build_testset())
tr_mae = accuracy.mae(tr_pred)
tr_predicted_ratings_N = np.asarray([p.est for p in tr_pred], dtype=np.float64)
print(tr_predicted_ratings_N[:10])

## Measure predictions on test set
print("Making predictions on test set (showing first 10):")
te_pred = model.test(test_set)
#te_mae = accuracy.mae(te_pred) # should be NaN because no real labels on testset
te_predicted_ratings_N = np.asarray([p.est for p in te_pred], dtype=np.float64)
print(te_predicted_ratings_N[:10])

print("n_factors %d tr_MAE %7.3f test_MAE %7.3f" % (n_factors, tr_mae, te_mae))

#print("Making test set predictions in the original order")
tep4=[]
for row in test_df.values[:]:
    userid = row[0]
    itemid = row[1]
    rhat = model.predict(userid, itemid)
    #print("user %4d item %4d predicted rating % 8.3f" % (userid, itemid, rhat.est))
    tep4.append(rhat.est)
np.savetxt('predicted_ratings_test5Mp.txt', np.asarray(tep4))

```

Computing the pearson similarity matrix...

Done computing similarity matrix.

Making predictions on training set (showing first 10):

MAE: 0.5375

```
[2.52974265 3.7955393 3.02586446 4.18648324 2.9166546 3.13912557
3.62204029 3.24213867 3.80484038 2.38314212]
```

Making predictions on test set (showing first 10):

```
[3.5294804 3.5294804 3.5294804 3.5294804 3.5294804 3.5294804 3.5294804
3.5294804 3.5294804 3.5294804]
```

```
n_factors      2 tr_MAE    0.538 test_MAE      nan
```

I use surprise KNN for problem 5. I look into the traditional KNN, KNN-Means(KNNM, subtract means), KNN using pearson correlation(KNNP) and KNN-Means with pearson correlation(KNNMP). I looked into how k affects the error of the methods. It seems choosing K=20 should be a good strategy in all of the cases. We observe the test error behavior as the following:

KNN: 0.7653

KNNM: 0.7427

KNNP: 0.7953

KNNMP: 0.7367

This shows when using KNN here, subtract mean and use pearson correlation is better for prediction.

In []:

